

```

1 from google.colab import drive
2 drive.mount('/content/drive')

→ Mounted at /content/drive

1
2 import pandas as pd
3 csv_path = './content/MyDrive/drive/data/csv'
4 artist_train = pd.read_csv('/content/drive/MyDrive/data/csv/Artist/
    artist_train')
5 # lets visualize one imag
6 base_url = '/content/drive/MyDrive/data/images'

1 # lets start creating data
2 artist = '/content/drive/MyDrive/data/csv/Artist'
3 genre = '/content/drive/MyDrive/data/csv/Genre'
4 style = '/content/drive/MyDrive/data/csv/Style'
5 data_dir = '/content/drive/MyDrive/data/csv'
6
7 artist_train_path = data_dir + '/artist_train.csv'
8 artist_val_path = data_dir + '/artist_val.csv'
9 artist_class_path = data_dir + '/artist_class.txt'
10
11 genre_train_path = data_dir + '/genre_train.csv'
12 genre_val_path = data_dir + '/genre_val.csv'
13 genre_class_path = data_dir + '/genre_class.txt'
14
15 style_train_path = data_dir + '/style_train.csv'
16 style_val_path = data_dir + '/style_val.csv'
17 style_class_path = data_dir + '/style_class.txt'
18
19 artist_train = pd.read_csv(data_dir + '/artist_train.csv')
20 artist_val = pd.read_csv(data_dir + '/artist_val.csv')
21 artist_class = pd.read_csv(artist_class_path, header=None, names=
    ["artist_name"])
22
23 genre_train = pd.read_csv(data_dir + '/genre_train.csv')
24 genre_val = pd.read_csv(data_dir + '/genre_val.csv')
25 genre_class = pd.read_csv(genre_class_path, header=None, names=["genre_name"])
26
27
28 style_train = pd.read_csv(data_dir + '/style_train.csv')
29 style_val = pd.read_csv(data_dir + '/style_val.csv')
30 style_class = pd.read_csv(style_class_path, header=None, names=["style_name"])
31
32 # genre_class['genre_name'][1]
33 len(style_class)

```

→ 27

```

1 import os
2 import pandas as pd
3 import torch
4 from torch.utils.data import Dataset
5 from torchvision import transforms
6 from PIL import Image
7 from collections import defaultdict
8 import random
9 from tqdm import tqdm
10 import matplotlib.pyplot as plt
11
12 # Define dataset class
13 class BalancedArtDataset(Dataset):
14     def __init__(self, csv_file, img_dir, class_mapping, transform=None, images_per_class=32):
15         self.data = pd.read_csv(csv_file)
16         self.img_dir = img_dir
17         self.class_mapping = class_mapping
18         self.transform = transform
19         self.images_per_class = images_per_class
20
21     # Filter out missing images
22     print("Filtering missing images...")
23     self.data = self.data[self.data.iloc[:, 0].apply(lambda x: os.path.exists(os.path.join(img_dir, str(x))))]
24

```

```

25     # Group images by class
26     print("Grouping images by class...")
27     self.class_images = defaultdict(list)
28     for _, row in tqdm(self.data.iterrows(), total=len(self.data), desc="Processing rows"):
29         self.class_images[row.iloc[1]].append(row)
30
31     # Balance dataset with 32 images per class
32     print("Balancing dataset...")
33     self.final_data = []
34     all_images = []
35     for cls, images in tqdm(self.class_images.items(), total=len(self.class_images), desc="Processing classes"):
36         if len(images) >= images_per_class:
37             selected_images = random.sample(images, images_per_class)
38         else:
39             selected_images = images[:]
40             all_images.extend(images) # Store extra images for filling
41         self.final_data.extend(selected_images)
42
43     # Fill missing slots with extra images
44     print("Filling missing slots...")
45     needed_images = images_per_class * len(self.class_images) - len(self.final_data)
46     if needed_images > 0:
47         self.final_data.extend(random.sample(all_images, min(needed_images, len(all_images))))
48
49     # Shuffle dataset
50     print("Shuffling dataset...")
51     random.shuffle(self.final_data)
52
53     # Count images per class
54     self.class_counts = defaultdict(int)
55     for row in self.final_data:
56         self.class_counts[row.iloc[1]] += 1
57
58     def __len__(self):
59         return len(self.final_data)
60
61     def __getitem__(self, idx):
62         row = self.final_data[idx]
63         img_path = os.path.join(self.img_dir, str(row.iloc[0]))
64         label = row.iloc[1]
65         image = Image.open(img_path).convert("RGB")
66
67         if self.transform:
68             image = self.transform(image)
69
70         return image, label, img_path
71
72
73     def visualize_class_distribution(self):
74         plt.figure(figsize=(12, 6))
75         plt.bar(self.class_counts.keys(), self.class_counts.values(), color='skyblue')
76         plt.xlabel("Class")
77         plt.ylabel("Number of Images")
78         plt.title("Class Distribution in Balanced Dataset")
79         plt.xticks(rotation=45)
80         plt.show()
81
82     def visualize_samples(self, num_samples=10):
83         fig, axes = plt.subplots(1, num_samples, figsize=(40, 20))
84         for i in range(num_samples):
85             image, label, img_path = self.__getitem__(random.randint(0, len(self) - 1))
86             image = image.permute(1, 2, 0).numpy() # Convert to (H, W, C)
87             image = (image * 0.5) + 0.5 # Unnormalize
88             axes[i].imshow(image)
89             axes[i].set_title(f"Class: {label} {artist_class['artist_name'][label]}")
90             axes[i].axis("off")
91         plt.show()
92
93     # Function to compare artist and genre relationships
94
95
96
97 # Define transformations
98 transform = transforms.Compose([
99     transforms.Resize((224, 224)),
100    transforms.ToTensor(),
101    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])

```

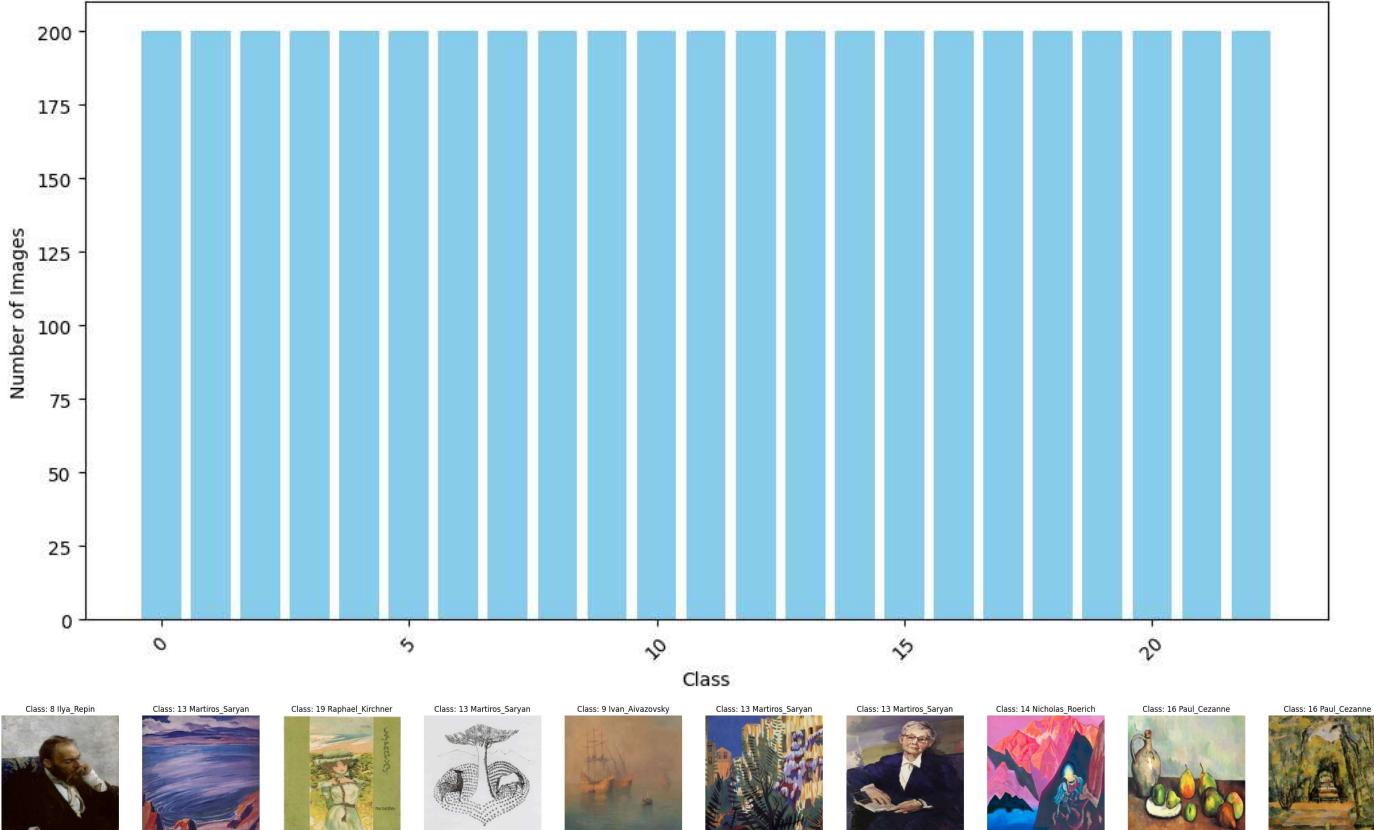
```

102 ])
103
104 # Create balanced artist dataset
105 print("Creating balanced artist dataset...")
106 artist_balanced_dataset = BalancedArtDataset(artist_train_path, "/content/drive/MyDrive/data/images", artist_class_path, transform=transform)
107 artist_test_balanced_dataset = BalancedArtDataset(artist_val_path, "/content/drive/MyDrive/data/images", artist_class_path, transform=transform)
108
109 # Check dataset length
110 print("Artist balanced dataset size:", len(artist_balanced_dataset), len(artist_test_balanced_dataset))
111
112 # Visualize class distribution
113 artist_balanced_dataset.visualize_class_distribution()
114
115 # Visualize sample images
116 artist_balanced_dataset.visualize_samples(num_samples=10)
117

```

↳ Creating balanced artist dataset...
 Filtering missing images...
 Grouping images by class...
 Processing rows: 100%|██████████| 13344/13344 [00:00<00:00, 14836.35it/s]
 Balancing dataset...
 Processing classes: 100%|██████████| 23/23 [00:00<00:00, 6719.77it/s]
 Filling missing slots...
 Shuffling dataset...
 Filtering missing images...
 Grouping images by class...
 Processing rows: 100%|██████████| 5706/5706 [00:00<00:00, 22442.96it/s]
 Balancing dataset...
 Processing classes: 100%|██████████| 23/23 [00:00<00:00, 19878.22it/s]
 Filling missing slots...
 Shuffling dataset...
 Artist balanced dataset size: 4600 1150

Class Distribution in Balanced Dataset



```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision.models as models
5 from torch.utils.data import DataLoader
6 from tqdm import tqdm
7
8 # Check for GPU
9 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

```

```
10 print(f"Using device: {device}")
11
→ Using device: cuda

1 # Define DataLoaders
2 batch_size = 32 # Adjust based on GPU memory
3
4 train_loader = DataLoader(artist_balanced_dataset, batch_size=batch_size, shuffle=True, num_workers=4,pin_memory=True)
5 val_loader = DataLoader(artist_test_balanced_dataset, batch_size=batch_size, shuffle=False, num_workers=4, pin_memory=True)
6
7 print("Dataloaders created successfully!")
8

→ Dataloaders created successfully!
/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
warnings.warn(
← ━━━━━━ →
```

```
1 # Hyperparameters
2 num_epochs_fc = 5 # Train FC layer only for 5 epochs
3 num_epochs_finetune = 6 # Fine-tune entire model for 10 more epochs
4 initial_lr = 0.001 # Learning rate for training FC layer
5 finetune_lr = 1e-4 # Lower learning rate for fine-tuning
6 batch_size = 32
7 weight_decay = 1e-4 # L2 regularization
8

9 import torch
10 import torch.nn as nn
11 import torch.optim as optim
12 from torch.optim.lr_scheduler import ReduceLROnPlateau
13 from torchvision import datasets, transforms, models
14 from torch.utils.data import DataLoader
15 from tqdm import tqdm
16 import time
17 # Load pre-trained ResNet-18 model
18 # Load pre-trained ResNet-18 model
19 num_classes = len(artist_class) # Number of output classes
20 model = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
21 model.fc = nn.Linear(model.fc.in_features, num_classes).to(device) # Replace final FC layer
22 model = model.to(device)
23
24 # Define loss function and optimizer for FC layer training
25 criterion = nn.CrossEntropyLoss()
26 fc_optimizer = optim.Adam(model.fc.parameters(), lr=initial_lr, weight_decay=weight_decay)
27
28 # Function to train the FC layer only
29 def train_fc_layer():
30     print("Starting Step 1: Training FC Layer Only...")
31
32     # Freeze all layers except the FC layer
33     for param in model.parameters():
34         param.requires_grad = False # Freeze all pretrained layers
35
36     # Ensure the new FC layer has requires_grad=True
37     for param in model.fc.parameters():
38         param.requires_grad = True
39
40     model.train() # Set the model to training mode
41
42     for epoch in range(num_epochs_fc):
43         running_loss, correct, total = 0.0, 0, 0
44
45         loop = tqdm(train_loader, leave=True, desc=f"Epoch [{epoch+1}/{num_epochs_fc}]")
46         for images, labels in loop:
47             images, labels = images.to(device), labels.to(device)
48             fc_optimizer.zero_grad() # Zero the gradients
49             outputs = model(images) # Forward pass
50             loss = criterion(outputs, labels) # Compute loss
51             loss.backward() # Backward pass to compute gradients
52             fc_optimizer.step() # Update FC layer weights
53
54             running_loss += loss.item()
55             _, predicted = torch.max(outputs, 1)
```

```

48     correct += (predicted == labels).sum().item()
49     total += labels.size(0)
50     loop.set_postfix(loss=running_loss / len(train_loader), acc=100 * correct / total)
51
52     print(f"Epoch [{epoch+1}/{num_epochs_fc}], Loss: {running_loss / len(train_loader):.4f}, "
53           f"Accuracy: {100 * correct / total:.2f}%")
54
55     print("⌚ Step 1 Complete: FC Layer Training Finished!")
56
57 # Step 2: Fine-Tuning Entire Network
58 def fine_tune():
59     print("Starting Step 2: Fine-Tuning the Entire Network...")
60     for param in model.parameters():
61         param.requires_grad = True # Unfreeze all layers
62
63     # Define optimizer and scheduler for fine-tuning
64     finetune_optimizer = optim.Adam(model.parameters(), lr=finetune_lr, weight_decay=weight_decay)
65     scheduler = ReduceLROnPlateau(finetune_optimizer, mode='min', factor=0.1, patience=3, verbose=True)
66
67     best_val_loss = float('inf')
68
69     for epoch in range(num_epochs_finetune):
70         model.train()
71         running_loss, correct, total = 0.0, 0, 0
72         loop = tqdm(train_loader, leave=True, desc=f"Epoch [{epoch+1}/{num_epochs_finetune}]")
73
74         for images, labels in loop:
75             images, labels = images.to(device), labels.to(device)
76             finetune_optimizer.zero_grad()
77             outputs = model(images)
78             loss = criterion(outputs, labels)
79             loss.backward()
80             finetune_optimizer.step()
81
82             running_loss += loss.item()
83             _, predicted = torch.max(outputs, 1)
84             correct += (predicted == labels).sum().item()
85             total += labels.size(0)
86             loop.set_postfix(loss=running_loss / len(train_loader), acc=100 * correct / total)
87
88     # Validation phase
89     val_loss, val_acc = evaluate(val_loader)
90     scheduler.step(val_loss)
91
92     print(f"Epoch [{epoch+1}/{num_epochs_finetune}], "
93           f"Train Loss: {running_loss / len(train_loader):.4f}, Train Acc: {100 * correct / total:.2f}%, "
94           f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")
95
96     # Save the best model based on validation loss
97     if val_loss < best_val_loss:
98         best_val_loss = val_loss
99         torch.save(model.state_dict(), "/content/drive/MyDrive/art_model/best_resnet18_model_200.pth")
100        print("✅ Best model saved!")
101
102    print("⌚ Step 2 Complete: Fine-Tuning Finished!")
103
104 # Function to evaluate the model on validation data
105 def evaluate(loader):
106     model.eval()
107     val_loss, correct, total = 0.0, 0, 0
108
109     with torch.no_grad():
110         for images, labels in loader:
111             images, labels = images.to(device), labels.to(device)
112             outputs = model(images)
113             loss = criterion(outputs, labels)
114             val_loss += loss.item()
115             _, predicted = torch.max(outputs, 1)
116             correct += (predicted == labels).sum().item()
117             total += labels.size(0)
118
119     avg_loss = val_loss / len(loader)
120     accuracy = 100 * correct / total
121     return avg_loss, accuracy
122
123
124 # Start Training: Step 1 (Train FC Layer) and Step 2 (Fine-Tuning)

```

```

125 train_fc_layer() # Train the FC layer only
126 fine_tune() # Fine-tune the entire ResNet model
127
128 print("🎉 Training complete! The best model is saved as 'best_resnet18_model.pth'.")

→ Starting Step 1: Training FC Layer Only...
Epoch [1/5]: 100%|██████████| 144/144 [04:22<00:00,  1.83s/it, acc=34.1, loss=2.38]
Epoch [1/5], Loss: 2.3813, Accuracy: 34.09%
Epoch [2/5]: 100%|██████████| 144/144 [01:17<00:00,  1.86it/s, acc=54.7, loss=1.64]
Epoch [2/5], Loss: 1.6399, Accuracy: 54.67%
Epoch [3/5]: 100%|██████████| 144/144 [01:15<00:00,  1.90it/s, acc=60.5, loss=1.4]
Epoch [3/5], Loss: 1.3960, Accuracy: 60.50%
Epoch [4/5]: 100%|██████████| 144/144 [01:17<00:00,  1.87it/s, acc=64.5, loss=1.25]
Epoch [4/5], Loss: 1.2537, Accuracy: 64.48%
Epoch [5/5]: 100%|██████████| 144/144 [01:15<00:00,  1.91it/s, acc=66.5, loss=1.16]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get
    warnings.warn(
Epoch [5/5], Loss: 1.1638, Accuracy: 66.46%
👉 Step 1 Complete: FC Layer Training Finished!
Starting Step 2: Fine-Tuning the Entire Network...
Epoch [1/6]: 100%|██████████| 144/144 [01:19<00:00,  1.81it/s, acc=71, loss=0.98]
Epoch [1/6], Train Loss: 0.9796, Train Acc: 71.00%, Val Loss: 0.9080, Val Acc: 73.13%
✓ Best model saved!
Epoch [2/6]: 100%|██████████| 144/144 [01:20<00:00,  1.80it/s, acc=96.9, loss=0.188]
Epoch [2/6], Train Loss: 0.1881, Train Acc: 96.93%, Val Loss: 0.8028, Val Acc: 76.26%
✓ Best model saved!
Epoch [3/6]: 100%|██████████| 144/144 [01:19<00:00,  1.80it/s, acc=99.8, loss=0.0509]
Epoch [3/6], Train Loss: 0.0509, Train Acc: 99.83%, Val Loss: 0.7850, Val Acc: 77.30%
✓ Best model saved!
Epoch [4/6]: 100%|██████████| 144/144 [01:20<00:00,  1.79it/s, acc=100, loss=0.0227]
Epoch [4/6], Train Loss: 0.0227, Train Acc: 100.00%, Val Loss: 0.7812, Val Acc: 78.17%
✓ Best model saved!
Epoch [5/6]: 100%|██████████| 144/144 [01:19<00:00,  1.81it/s, acc=100, loss=0.0137]
Epoch [5/6], Train Loss: 0.0137, Train Acc: 100.00%, Val Loss: 0.7947, Val Acc: 77.65%
Epoch [6/6]: 100%|██████████| 144/144 [01:19<00:00,  1.82it/s, acc=100, loss=0.00997]
Epoch [6/6], Train Loss: 0.0100, Train Acc: 100.00%, Val Loss: 0.7785, Val Acc: 78.52%
✓ Best model saved!
👉 Step 2 Complete: Fine-Tuning Finished!
🎉 Training complete! The best model is saved as 'best_resnet18_model.pth'.

```

```

1 model.load_state_dict(torch.load("/content/drive/MyDrive/art_model/best_resnet18_model_200.pth", map_location=device))
2 model = model.to(device)
3
4 # Set the model to evaluation mode for inference
5 model.eval()
6

```

```

→ ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (1): BasicBlock(

```

```
(conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
(conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(layer3): Sequential(
(0): BasicBlock(
(conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(downsample): Sequential(
(0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
(1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(1): BasicBlock(
(conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```

1

```
1 import torch
2
3 # Load the entire model, explicitly setting `weights_only=False`
4 loaded_model = torch.load('/content/drive/MyDrive/art_model/artist_model_02_resnet_18.pth', weights_only=False)
5 loaded_model
```

ResNet(
 (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
 (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (relu): ReLU(inplace=True)
 (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
 (layer1): Sequential(
 (0): BasicBlock(
 (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (relu): ReLU(inplace=True)
 (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
 (1): BasicBlock(
 (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (relu): ReLU(inplace=True)
 (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
 (layer2): Sequential(
 (0): BasicBlock(
 (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
 (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (relu): ReLU(inplace=True)
 (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (downsample): Sequential(
 (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
 (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
 (1): BasicBlock(
 (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (relu): ReLU(inplace=True)
 (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
 (layer3): Sequential(
 (0): BasicBlock(
 (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
 (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (relu): ReLU(inplace=True)
 (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
 (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 (downsample): Sequential(
 (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
)

```

        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)

1 import torch
2 import matplotlib.pyplot as plt
3 import random
4 random.seed(42)
5
6 # Function to visualize random predictions from test dataset
7 def visualize_random_predictions(model, dataset, class_mapping, num_samples=8):
8     model.eval() # Set model to evaluation mode
9     fig, axes = plt.subplots(1, num_samples, figsize=(20, 8))
10
11    with torch.no_grad(): # Disable gradient calculation for inference
12        for i in range(num_samples):
13            # Pick a random index from the dataset
14            idx = random.randint(0, len(dataset) - 1)
15            image, true_label = dataset[idx]
16
17            # Add batch dimension and move image to the appropriate device
18            image_batch = image.unsqueeze(0).to(device)
19
20            # Make a prediction
21            output = model(image_batch)
22            _, predicted_label = torch.max(output, 1)
23
24            # Convert image tensor to numpy and unnormalize
25            image = image.permute(1, 2, 0).cpu().numpy() # Convert (C, H, W) -> (H, W, C)
26            image = (image * 0.5) + 0.5 # Unnormalize image
27
28            # Display the image and predictions
29            axes[i].imshow(image)
30            axes[i].set_title(
31                f"True: {class_mapping['artist_name'][true_label]}\nPredicted: {class_mapping['artist_name'][predicted_label.item()]}", color="green" if true_label == predicted_label.item() else "red"), font-size=12,
32            )
33            axes[i].axis("off")
34
35    plt.show()
36
37
38
39
40 # Visualize random predictions on test dataset (Adjust `num_samples` as needed)
41 visualize_random_predictions(model, artist_test_balanced_dataset, artist_class, num_samples=8)
42 # Visualize random predictions on test dataset (Adjust `num_samples` as needed)
43 visualize_random_predictions(loaded_model, artist_test_balanced_dataset, artist_class, num_samples=8)
44

```



```

1 import torch
2 import matplotlib.pyplot as plt
3 import random
4
5
6 # Function to visualize random predictions from test dataset on the same images for both models

```

```

7 def visualize_same_random_predictions(model1, model2, dataset, class_mapping, num_samples=8):
8     model1.eval() # Set models to evaluation mode
9     model2.eval()
10
11    fig, axes = plt.subplots(2, num_samples, figsize=(20, 12)) # Two rows: one for each model
12
13    # Generate fixed random indices to ensure both models predict on the same images
14    fixed_indices = [random.randint(0, len(dataset) - 1) for _ in range(num_samples)]
15
16    with torch.no_grad(): # Disable gradient calculation for inference
17        for i, idx in enumerate(fixed_indices):
18            image, true_label = dataset[idx]
19
20            # Add batch dimension and move image to the appropriate device
21            image_batch = image.unsqueeze(0).to(device)
22
23            # Make predictions with both models
24            output1 = model1(image_batch)
25            output2 = model2(image_batch)
26            _, predicted_label1 = torch.max(output1, 1)
27            _, predicted_label2 = torch.max(output2, 1)
28
29            # Convert image tensor to numpy and unnormalize
30            image = image.permute(1, 2, 0).cpu().numpy() # Convert (C, H, W) -> (H, W, C)
31            image = (image * 0.5) + 0.5 # Unnormalize image
32
33            # Display the image and predictions for model1 (top row)
34            axes[0, i].imshow(image)
35            axes[0, i].set_title(
36                f"Model 1 - True: {class_mapping['artist_name'][true_label]}\nPredicted: {class_mapping['artist_name'][predicted_label1]}"
37                color="green" if true_label == predicted_label1.item() else "red",
38                fontsize=12,
39            )
40            axes[0, i].axis("off")
41
42            # Display the image and predictions for model2 (bottom row)
43            axes[1, i].imshow(image)
44            axes[1, i].set_title(
45                f"Model 2 - True: {class_mapping['artist_name'][true_label]}\nPredicted: {class_mapping['artist_name'][predicted_label2]}"
46                color="green" if true_label == predicted_label2.item() else "red",
47                fontsize=12,
48            )
49            axes[1, i].axis("off")
50
51    plt.show()
52
53
54 # Visualize random predictions on test dataset for both models on the same set of images
55 visualize_same_random_predictions(model, loaded_model, artist_test_balanced_dataset, artist_class, num_samples=8)
56

```



Model 1 - True: Albrecht_Durer
 Model 1 - True: Raphael_Kirchner
 Model 1 - True: Edgar_Degas
 Model 1 - True: Vincent_van_Gogh
 Model 1 - True: Pablo_Picasso
 Model 1 - True: Albrecht_Durer
 Model 1 - True: Marc_Chagall
 Model 1 - True: Eugene_Boudin
 Predicted: Albrecht_Durer
 Predicted: Raphael_Kirchner
 Predicted: Edgar_Degas
 Predicted: Vincent_van_Gogh
 Predicted: Pablo_Picasso
 Predicted: Ivan_Shishkin
 Predicted: Marc_Chagall
 Predicted: Eugene_Boudin



Model 2 - True: Albrecht_Durer
 Model 2 - True: Raphael_Kirchner
 Model 2 - True: Edgar_Degas
 Model 2 - True: Vincent_van_Gogh
 Model 2 - True: Pablo_Picasso
 Model 2 - True: Albrecht_Durer
 Model 2 - True: Marc_Chagall
 Model 2 - True: Eugene_Boudin
 Predicted: Albrecht_Durer
 Predicted: Albrecht_Durer
 Predicted: Edgar_Degas
 Predicted: Vincent_van_Gogh
 Predicted: Pablo_Picasso
 Predicted: Childe_Hassam
 Predicted: Pablo_Picasso
 Predicted: Childe_Hassam



```

1 import torch
2 import matplotlib.pyplot as plt
3 import random
4
5 # Function to visualize random predictions with true artist, style, and genre classes
6 def visualize_predictions_with_true_classes(model1, model2, dataset, artist_class, style_class, genre_class, style_val, genre_val, num_samples):
7     model1.eval() # Set models to evaluation mode
8     model2.eval()
9
10    # Normalize paths in style_val and genre_val
11    style_val2 = {
12        path.strip().replace("\\", "/"): class_no
13        for path, class_no in zip(style_val['path'], style_val['class_no'])}
14    }
15    genre_val2 = {
16        path.strip().replace("\\", "/"): class_no
17        for path, class_no in zip(genre_val['path'], genre_val['class_no'])}
18    }
19
20    fig, axes = plt.subplots(2, num_samples, figsize=(20, 12)) # Two rows: one for each model
21
22    # Generate fixed random indices to ensure both models predict on the same images
23    fixed_indices = [random.randint(0, len(dataset) - 1) for _ in range(num_samples)]
24
25    with torch.no_grad(): # Disable gradient calculation for inference
26        for i, idx in enumerate(fixed_indices):
27            # Unpack image, true artist label, and image path
28            image, true_artist_label, image_path = dataset[idx]
29
30            print(f"Original Image Path: {image_path}") # Debugging line
31
32            # Modify the image path to extract the relative path after "images/"
33            relative_path = image_path.split("images/")[-1] # Extract relative path
34            relative_path = relative_path.strip().replace("\\", "/") # Normalize path
35
36            print(f"Relative Path: {relative_path}") # Debugging line
37
38            # Check if the relative path exists in the mappings, and handle missing cases
39            if relative_path in style_val2 and relative_path in genre_val2:
40                true_style_label = style_val2[relative_path] # Use the normalized dictionary
41                true_genre_label = genre_val2[relative_path] # Use the normalized dictionary
42            else:
43                print(f"Warning: Relative path {relative_path} not found in style_val or genre_val.")
44                continue # Skip the sample if not found in the mappings
45
46            # Add batch dimension and move image to the appropriate device
47            image_batch = image.unsqueeze(0).to(device)
48
49            # Make predictions with both models
50            output1 = model1(image_batch)
51            output2 = model2(image_batch)
52
53            # Get predicted artist labels for both models
54            _, predicted_artist1 = torch.max(output1, 1)
55            _, predicted_artist2 = torch.max(output2, 1)
56
57            # Convert image tensor to numpy and unnormalize
58            image = image.permute(1, 2, 0).cpu().numpy() # Convert (C, H, W) -> (H, W, C)
59            image = (image * 0.5) + 0.5 # Unnormalize image
60
61            # Display the image and predictions for model1 (top row)
62            axes[0, i].imshow(image)
63            axes[0, i].set_title(
64                f"Model 1\nArtist: {artist_class['artist_name'][predicted_artist1.item()]}\n"
65                f"True Style: {style_class['style_name'][true_style_label]}\n"
66                f"True Genre: {genre_class['genre_name'][true_genre_label]},\n"
67                f"color='green' if true_artist_label == predicted_artist1.item() else 'red',\n"
68                f"fontsize=10,\n"
69            )
70            axes[0, i].axis("off")
71
72            # Display the image and predictions for model2 (bottom row)
73            axes[1, i].imshow(image)
74            axes[1, i].set_title(
75                f"Model 2\nArtist: {artist_class['artist_name'][predicted_artist2.item()]}\n"
76                f"True Style: {style_class['style_name'][true_style_label]}\n"

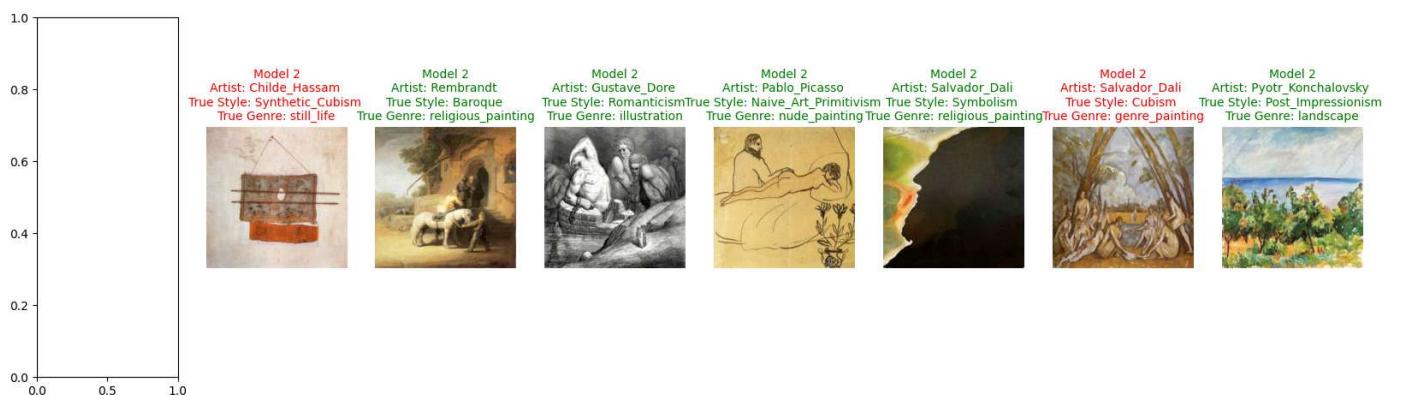
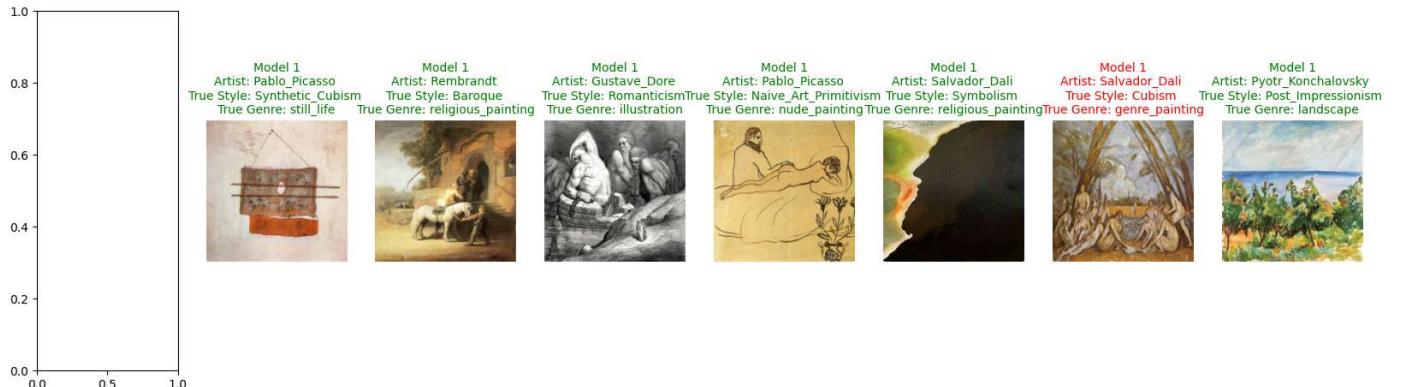
```

```

77         f"True Genre: {genre_class['genre_name'][true_genre_label]}",
78         color=("green" if true_artist_label == predicted_artist2.item() else "red"),
79         fontsize=10,
80     )
81     axes[1, i].axis("off")
82
83 plt.show()
84
85 # Example usage of the function (assuming `style_val` and `genre_val` are defined mappings)
86 visualize_predictions_with_true_classes(
87     model,
88     loaded_model,
89     artist_test_balanced_dataset,
90     artist_class,
91     style_class,
92     genre_class,
93     style_val, # Dictionary mapping relative image paths to style class indices
94     genre_val, # Dictionary mapping relative image paths to genre class indices
95     num_samples=8
96 )
97

```

Original Image Path: /content/drive/MyDrive/data/images/Art_Nouveau_Modern/raphael-kirchner_roma-6.jpg
 Relative Path: Art_Nouveau_Modern/raphael-kirchner_roma-6.jpg
 Warning: Relative path Art_Nouveau_Modern/raphael-kirchner_roma-6.jpg not found in style_val or genre_val.
 Original Image Path: /content/drive/MyDrive/data/images/Synthetic_Cubism/pablo-picasso_guitar.jpg
 Relative Path: Synthetic_Cubism/pablo-picasso_guitar.jpg
 Original Image Path: /content/drive/MyDrive/data/images/Baroque/rembrandt_charitable-samaritan-also-known-as-the-good-samaritan-1638.jpg
 Relative Path: Baroque/rembrandt_charitable-samaritan-also-known-as-the-good-samaritan-1638.jpg
 Original Image Path: /content/drive/MyDrive/data/images/Romanticism/gustave-dore_the-inferno-canto-31.jpg
 Relative Path: Romanticism/gustave-dore_the-inferno-canto-31.jpg
 Original Image Path: /content/drive/MyDrive/data/images/Naive_Art_Primitivism/pablo-picasso_nude-with-picasso-by-her-feet.jpg
 Relative Path: Naive_Art_Primitivism/pablo-picasso_nude-with-picasso-by-her-feet.jpg
 Original Image Path: /content/drive/MyDrive/data/images/Symbolism/salvador-dali_abraham-pater-multarem-gentium-genesis-12-1f-1967.jpg
 Relative Path: Symbolism/salvador-dali_abraham-pater-multarem-gentium-genesis-12-1f-1967.jpg
 Original Image Path: /content/drive/MyDrive/data/images/Cubism/paul-cezanne_large-bathers-1906-1.jpg
 Relative Path: Cubism/paul-cezanne_large-bathers-1906-1.jpg
 Original Image Path: /content/drive/MyDrive/data/images/Post_Impressionism/pyotr-konchalovsky_gurzuf-the-trees-on-the-background-of-the-sea-1929.jpg
 Relative Path: Post_Impressionism/pyotr-konchalovsky_gurzuf-the-trees-on-the-background-of-the-sea-1929.jpg



```

1 # Load your DataFrame (replace this with the path to your actual CSV)
2 # style_val_df = pd.read_csv('style_values.csv')

```

```

3
4 # Convert DataFrame to dictionary with normalized paths
5 style_val2 = {
6     path.strip().replace("\\", "/"): class_no
7     for path, class_no in zip(style_val['path'], style_val['class_no'])
8 }
9 key = 'Impressionism/edgar-degas_dancers-on-set-1880.jpg' # Your specific path
10 key = key.strip().replace("\\", "/") # Normalize the path to match dictionary keys
11
12 if key in style_val2:
13     print("Found! Class number:", style_val2[key])
14 else:
15     print(f"Key '{key}' not found in style_val2.")
16

```

↳ Found! Class number: 12

1 style_val

	path	class_no
0	Impressionism/edgar-degas_dancers-on-set-1880.jpg	12
1	Impressionism/claudie-monet_water-lilies-6.jpg	12
2	Impressionism/giovanni-boldini_a-guitar-player...	12
3	Impressionism/john-singer-sargent_at-torre-gal...	12
4	Impressionism/john-singer-sargent_artist-in-th...	12
...
24416	Naive_Art_Primitivism/niko-pirosmani_leaning-a...	15
24417	Naive_Art_Primitivism/marc-chagall_russian-wed...	15
24418	Naive_Art_Primitivism/camille-bombois_port-de-...	15
24419	Naive_Art_Primitivism/marc-chagall_god-directs...	15
24420	Naive_Art_Primitivism/natalia-goncharova_sheep...	15

24421 rows × 2 columns

▼ Style dataset

```

1 import os
2 import pandas as pd
3 import torch
4 from torch.utils.data import Dataset
5 from torchvision import transforms
6 from PIL import Image
7 from collections import defaultdict
8 import random
9 from tqdm import tqdm
10 import matplotlib.pyplot as plt
11
12 # Define dataset class
13 class BalancedArtDataset(Dataset):
14     def __init__(self, csv_file, img_dir, class_mapping, return_path = False, transform=None, images_per_class=32):
15         self.data = pd.read_csv(csv_file)
16         self.img_dir = img_dir
17         self.class_mapping = class_mapping
18         self.transform = transform
19         self.images_per_class = images_per_class
20         self.return_path = return_path # Ensure this is set as an attribute
21
22
23     # Filter out missing images
24     print("Filtering missing images...")
25     self.data = self.data[self.data.iloc[:, 0].apply(lambda x: os.path.exists(os.path.join(img_dir, str(x))))]
26
27     # Group images by class
28     print("Grouping images by class...")
29     self.class_images = defaultdict(list)
30     for _, row in tqdm(self.data.iterrows(), total=len(self.data), desc="Processing rows"):
31         self.class_images[row.iloc[1]].append(row)

```

```

32
33     # Balance dataset with 32 images per class
34     print("Balancing dataset...")
35     self.final_data = []
36     all_images = []
37     for cls, images in tqdm(self.class_images.items(), total=len(self.class_images), desc="Processing classes"):
38         if len(images) >= images_per_class:
39             selected_images = random.sample(images, images_per_class)
40         else:
41             selected_images = images[:]
42             all_images.extend(images) # Store extra images for filling
43         self.final_data.extend(selected_images)
44
45     # Fill missing slots with extra images
46     print("Filling missing slots...")
47     needed_images = images_per_class * len(self.class_images) - len(self.final_data)
48     if needed_images > 0:
49         self.final_data.extend(random.sample(all_images, min(needed_images, len(all_images))))
50
51     # Shuffle dataset
52     print("Shuffling dataset...")
53     random.shuffle(self.final_data)
54
55     # Count images per class
56     self.class_counts = defaultdict(int)
57     for row in self.final_data:
58         self.class_counts[row.iloc[1]] += 1
59
60     def __len__(self):
61         return len(self.final_data)
62
63     def __getitem__(self, idx):
64         row = self.final_data[idx]
65         img_path = os.path.join(self.img_dir, str(row.iloc[0]))
66         label = row.iloc[1]
67         image = Image.open(img_path).convert("RGB")
68
69         if self.transform:
70             image = self.transform(image)
71
72         if self.return_path: # Return img_path if flag is set
73             return image, label, img_path
74         else: # Return only image and label for model training
75             return image, label
76
77
78     def visualize_class_distribution(self):
79         plt.figure(figsize=(12, 6))
80         plt.bar(self.class_counts.keys(), self.class_counts.values(), color='skyblue')
81         plt.xlabel("Class")
82         plt.ylabel("Number of Images")
83         plt.title("Class Distribution in Balanced Dataset")
84         plt.xticks(rotation=45)
85         plt.show()
86
87     def visualize_samples(self, num_samples=10):
88         fig, axes = plt.subplots(1, num_samples, figsize=(40, 20))
89         for i in range(num_samples):
90             image, label, img_path = self.__getitem__(random.randint(0, len(self) - 1))
91             image = image.permute(1, 2, 0).numpy() # Convert to (H, W, C)
92             image = (image * 0.5) + 0.5 # Unnormalize
93             axes[i].imshow(image)
94             axes[i].set_title(f"Class: {label} {artist_class['artist_name'][label]}")
95             axes[i].axis("off")
96         plt.show()
97
98     # Function to compare artist and genre relationships
99
100
101
102
103 # Define transformations
104 transform = transforms.Compose([
105     transforms.Resize((224, 224)),
106     transforms.ToTensor(),
107     transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
108 ])

```

```
109
110
```

```
1 style_balanced_dataset = BalancedArtDataset(style_train_path, "/content/drive/MyDrive/data/images", style_class_path, transform=transform
2 style_test_balanced_dataset = BalancedArtDataset(style_val_path, "/content/drive/MyDrive/data/images", style_class_path, transform=trans
3 len(style_balanced_dataset)
```

Filtering missing images...
 Grouping images by class...
 Processing rows: 100%|██████████| 55610/55610 [00:03<00:00, 17632.04it/s]
 Balancing dataset...
 Processing classes: 100%|██████████| 27/27 [00:00<00:00, 12056.45it/s]
 Filling missing slots...
 Shuffling dataset...
 Filtering missing images...
 Grouping images by class...
 Processing rows: 100%|██████████| 23818/23818 [00:01<00:00, 12965.21it/s]
 Balancing dataset...
 Processing classes: 100%|██████████| 27/27 [00:00<00:00, 13497.76it/s]Filling missing slots...
 Shuffling dataset...

```
5296
```

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision.models as models
5 from torch.utils.data import DataLoader
6 from tqdm import tqdm
7
8 # Check for GPU
9 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10 print(f"Using device: {device}")
11
```

Using device: cuda

```
1 # Define DataLoaders
2 batch_size = 32 # Adjust based on GPU memory
3
4 style_train_loader = DataLoader(style_balanced_dataset, batch_size=batch_size, shuffle=True, num_workers=4,pin_memory=True)
5 style_val_loader = DataLoader(style_test_balanced_dataset, batch_size=batch_size, shuffle=False, num_workers=4, pin_memory=True)
6
7 print("Dataloaders created successfully!")
8
9 # Hyperparameters
10 num_epochs_fc = 5 # Train FC layer only for 5 epochs
11 num_epochs_finetune = 5 # Fine-tune entire model for 10 more epochs
12 initial_lr = 0.001 # Learning rate for training FC layer
13 finetune_lr = 1e-4 # Lower learning rate for fine-tuning
14 batch_size = 32
15 weight_decay = 1e-4 # L2 regularization
16
```

Dataloaders created successfully!
 /usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
 warnings.warn(

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.optim.lr_scheduler import ReduceLROnPlateau
5 from torchvision import datasets, transforms, models
6 from torch.utils.data import DataLoader
7 from tqdm import tqdm
8 import time
9 # Load pre-trained ResNet-18 model
10 # Load pre-trained ResNet-18 model
11 num_classes = len(style_class) # Number of output classes
12 model_style = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
13 model_style.fc = nn.Linear(model_style.fc.in_features, num_classes).to(device) # Replace final FC layer
14 model_style = model_style.to(device)
15
16 # Define loss function and optimizer for FC layer training
17 criterion = nn.CrossEntropyLoss()
```

```

18 fc_optimizer = optim.Adam(model_style.fc.parameters(), lr=initial_lr, weight_decay=weight_decay)
19
20 # Function to train the FC layer only
21 def train_fc_layer():
22     print("Starting Step 1: Training FC Layer Only...")
23
24     # Freeze all layers except the FC layer
25     for param in model_style.parameters():
26         param.requires_grad = False # Freeze all pretrained layers
27
28     # Ensure the new FC layer has requires_grad=True
29     for param in model_style.fc.parameters():
30         param.requires_grad = True
31
32     model_style.train() # Set the model to training mode
33
34     for epoch in range(num_epochs_fc):
35         running_loss, correct, total = 0.0, 0, 0
36
37         loop = tqdm(style_train_loader, leave=True, desc=f"Epoch [{epoch+1}/{num_epochs_fc}]")
38         for images, labels in loop:
39             images, labels = images.to(device), labels.to(device)
40             fc_optimizer.zero_grad() # Zero the gradients
41             outputs = model_style(images) # Forward pass
42             loss = criterion(outputs, labels) # Compute loss
43             loss.backward() # Backward pass to compute gradients
44             fc_optimizer.step() # Update FC layer weights
45
46             running_loss += loss.item()
47             _, predicted = torch.max(outputs, 1)
48             correct += (predicted == labels).sum().item()
49             total += labels.size(0)
50         loop.set_postfix(loss=running_loss / len(style_train_loader), acc=100 * correct / total)
51
52         print(f"Epoch [{epoch+1}/{num_epochs_fc}], Loss: {running_loss / len(style_train_loader):.4f}, "
53               f"Accuracy: {100 * correct / total:.2f}%")
54
55     print("⌚ Step 1 Complete: FC Layer Training Finished!")
56
57 # Step 2: Fine-Tuning Entire Network
58 def fine_tune():
59     print("Starting Step 2: Fine-Tuning the Entire Network...")
60     for param in model_style.parameters():
61         param.requires_grad = True # Unfreeze all layers
62
63     # Define optimizer and scheduler for fine-tuning
64     finetune_optimizer = optim.Adam(model_style.parameters(), lr=finetune_lr, weight_decay=weight_decay)
65     scheduler = ReduceLROnPlateau(finetune_optimizer, mode='min', factor=0.1, patience=3, verbose=True)
66
67     best_val_loss = float('inf')
68
69     for epoch in range(num_epochs_finetune):
70         model_style.train()
71         running_loss, correct, total = 0.0, 0, 0
72         loop = tqdm(style_train_loader, leave=True, desc=f"Epoch [{epoch+1}/{num_epochs_finetune}]")
73
74         for images, labels in loop:
75             images, labels = images.to(device), labels.to(device)
76             finetune_optimizer.zero_grad()
77             outputs = model_style(images)
78             loss = criterion(outputs, labels)
79             loss.backward()
80             finetune_optimizer.step()
81
82             running_loss += loss.item()
83             _, predicted = torch.max(outputs, 1)
84             correct += (predicted == labels).sum().item()
85             total += labels.size(0)
86             loop.set_postfix(loss=running_loss / len(style_train_loader), acc=100 * correct / total)
87
88     # Validation phase
89     val_loss, val_acc = evaluate(style_val_loader)
90     scheduler.step(val_loss)
91
92     print(f"Epoch [{epoch+1}/{num_epochs_finetune}], "
93           f"Train Loss: {running_loss / len(style_train_loader):.4f}, Train Acc: {100 * correct / total:.2f}%, "
94           f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")

```

```

95
96     # Save the best model based on validation loss
97     if val_loss < best_val_loss:
98         best_val_loss = val_loss
99         torch.save(model_style.state_dict(), "/content/drive/MyDrive/art_model/best_resnet18_style_model_200.pth")
100        print("✅ Best model saved!")
101
102    print("⌚ Step 2 Complete: Fine-Tuning Finished!")
103
104 # Function to evaluate the model on validation data
105 def evaluate(loader):
106     model_style.eval()
107     val_loss, correct, total = 0.0, 0, 0
108
109     with torch.no_grad():
110         for images, labels in loader:
111             images, labels , = images.to(device), labels.to(device)
112             outputs = model_style(images)
113             loss = criterion(outputs, labels)
114             val_loss += loss.item()
115             _, predicted = torch.max(outputs, 1)
116             correct += (predicted == labels).sum().item()
117             total += labels.size(0)
118
119     avg_loss = val_loss / len(loader)
120     accuracy = 100 * correct / total
121     return avg_loss, accuracy
122
123
124 # Start Training: Step 1 (Train FC Layer) and Step 2 (Fine-Tuning)
125 train_fc_layer() # Train the FC layer only
126 fine_tune()       # Fine-tune the entire ResNet model
127
128 print("🎉 Training complete! The best model is saved as 'best_resnet18_model.pth'.")
```

⌚ Starting Step 1: Training FC Layer Only...

Epoch [1/5]: 14%|██████| 23/166 [00:25<02:20, 1.02it/s, acc=6.51, loss=0.483]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [1/5]: 100%|██████████| 166/166 [01:48<00:00, 1.53it/s, acc=23.5, loss=2.69]
Epoch [1/5], Loss: 2.6946, Accuracy: 23.51
Epoch [2/5]: 64%|██████| 107/166 [01:03<00:24, 2.39it/s, acc=40.4, loss=1.32]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [2/5]: 100%|██████████| 166/166 [01:36<00:00, 1.72it/s, acc=41.5, loss=2.02]
Epoch [2/5], Loss: 2.0170, Accuracy: 41.47%
Epoch [3/5]: 8%|██| 14/166 [00:11<01:34, 1.61it/s, acc=47.5, loss=0.175]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [3/5]: 100%|██████████| 166/166 [01:34<00:00, 1.75it/s, acc=46.9, loss=1.81]
Epoch [3/5], Loss: 1.8091, Accuracy: 46.90%
Epoch [4/5]: 56%|██████| 93/166 [00:57<00:29, 2.47it/s, acc=50, loss=0.935]/usr/local/lib/python3.11/dist-packages/PIL/Image.py:34 warnings.warn(
Epoch [4/5]: 100%|██████████| 166/166 [01:38<00:00, 1.69it/s, acc=49.7, loss=1.68]
Epoch [4/5], Loss: 1.6815, Accuracy: 49.66%
Epoch [5/5]: 51%|████| 84/166 [00:49<00:33, 2.44it/s, acc=52.5, loss=0.802]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [5/5]: 100%|██████████| 166/166 [01:41<00:00, 1.63it/s, acc=51.5, loss=1.6]
Epoch [5/5], Loss: 1.6000, Accuracy: 51.53%
⌚ Step 1 Complete: FC Layer Training Finished!

⌚ Starting Step 2: Fine-Tuning the Entire Network...

Epoch [1/5]: 0%| | 0/166 [00:00<?, ?it/s]/usr/local/lib/python3.11/dist-packages/PIL/Image.py:3402: DecompressionBombWarning: warnings.warn(
Epoch [1/5]: 100%|██████████| 166/166 [02:05<00:00, 1.33it/s, acc=54.1, loss=1.47]
Epoch [1/5], Train Loss: 1.4744, Train Acc: 54.12%, Val Loss: 1.6001, Val Acc: 50.15%
✅ Best model saved!
Epoch [2/5]: 14%|██| 24/166 [00:17<01:35, 1.49it/s, acc=86.5, loss=0.0888]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [2/5]: 100%|██████████| 166/166 [01:44<00:00, 1.59it/s, acc=86.9, loss=0.556]
Epoch [2/5], Train Loss: 0.5557, Train Acc: 86.88%, Val Loss: 1.5541, Val Acc: 52.52%
✅ Best model saved!
Epoch [3/5]: 88%|███████| 146/166 [01:31<00:11, 1.67it/s, acc=98, loss=0.171]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [3/5]: 100%|██████████| 166/166 [01:41<00:00, 1.64it/s, acc=97.9, loss=0.194]
Epoch [3/5], Train Loss: 0.1942, Train Acc: 97.92%, Val Loss: 1.5494, Val Acc: 52.67%
✅ Best model saved!
Epoch [4/5]: 10%|██| 16/166 [00:13<01:18, 1.91it/s, acc=99.2, loss=0.00788]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [4/5]: 100%|██████████| 166/166 [01:39<00:00, 1.67it/s, acc=99.6, loss=0.075]
Epoch [4/5], Train Loss: 0.0750, Train Acc: 99.57%, Val Loss: 1.5867, Val Acc: 53.26%
Epoch [5/5]: 64%|██████| 106/166 [01:04<00:28, 2.07it/s, acc=99.7, loss=0.0277]/usr/local/lib/python3.11/dist-packages/PIL/Image.py: warnings.warn(
Epoch [5/5]: 100%|██████████| 166/166 [01:41<00:00, 1.64it/s, acc=99.7, loss=0.0415]
Epoch [5/5], Train Loss: 0.0415, Train Acc: 99.72%, Val Loss: 1.6110, Val Acc: 53.78%

⌚ Step 2 Complete: Fine-Tuning Finished!
🎉 Training complete! The best model is saved as 'best_resnet18_model.pth'.

Genre dataset

```
1 genre_balanced_dataset = BalancedArtDataset(genre_train_path, "/content/drive/MyDrive/data/images", genre_class_path, transform=transform
2 genre_test_balanced_dataset = BalancedArtDataset(genre_val_path, "/content/drive/MyDrive/data/images", genre_class_path, transform=transform
3 len(genre_balanced_dataset)
```

→ Filtering missing images...
Grouping images by class...
Processing rows: 100% [██████] 44393/44393 [00:02<00:00, 16196.42it/s]
Balancing dataset...
Processing classes: 100% [██████] 10/10 [00:00<00:00, 1698.37it/s]
Filling missing slots...
Shuffling dataset...
Filtering missing images...
Grouping images by class...
Processing rows: 100% [██████] 19005/19005 [00:01<00:00, 16650.14it/s]
Balancing dataset...
Processing classes: 100% [██████] 10/10 [00:00<00:00, 9754.20it/s] Filling missing slots...
Shuffling dataset...

4500

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision.models as models
5 from torch.utils.data import DataLoader
6 from tqdm import tqdm
7
8 # Check for GPU
9 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10 print(f"Using device: {device}")
11
```

→ Using device: cuda

```
1 # Define DataLoaders
2 batch_size = 32 # Adjust based on GPU memory
3
4 genre_train_loader = DataLoader(genre_balanced_dataset, batch_size=batch_size, shuffle=True, num_workers=4, pin_memory=True)
5 genre_val_loader = DataLoader(genre_test_balanced_dataset, batch_size=batch_size, shuffle=False, num_workers=4, pin_memory=True)
6
7 print("Dataloaders created successfully!")
8
9 # Hyperparameters
10 num_epochs_fc = 5 # Train FC layer only for 5 epochs
11 num_epochs_finetune = 5 # Fine-tune entire model for 10 more epochs
12 initial_lr = 0.001 # Learning rate for training FC layer
13 finetune_lr = 1e-4 # Lower learning rate for fine-tuning
14 batch_size = 32
15 weight_decay = 1e-4 # L2 regularization
16
```

→ Dataloaders created successfully!
/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
warnings.warn(

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.optim.lr_scheduler import ReduceLROnPlateau
5 from torchvision import datasets, transforms, models
6 from torch.utils.data import DataLoader
7 from tqdm import tqdm
8 import time
9 # Load pre-trained ResNet-18 model
10 # Load pre-trained ResNet-18 model
11 num_classes = len(genre_class) # Number of output classes
12 model_genre = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)
```

```

13 model_genre.fc = nn.Linear(model_genre.fc.in_features, num_classes).to(device) # Replace final FC layer
14 model_genre = model_genre.to(device)
15
16 # Define loss function and optimizer for FC layer training
17 criterion = nn.CrossEntropyLoss()
18 fc_optimizer = optim.Adam(model_genre.fc.parameters(), lr=initial_lr, weight_decay=weight_decay)
19
20 # Function to train the FC layer only
21 def train_fc_layer():
22     print("Starting Step 1: Training FC Layer Only...")
23
24     # Freeze all layers except the FC layer
25     for param in model_genre.parameters():
26         param.requires_grad = False # Freeze all pretrained layers
27
28     # Ensure the new FC layer has requires_grad=True
29     for param in model_genre.fc.parameters():
30         param.requires_grad = True
31
32     model_genre.train() # Set the model to training mode
33
34     for epoch in range(num_epochs_fc):
35         running_loss, correct, total = 0.0, 0, 0
36
37     loop = tqdm(genre_train_loader, leave=True, desc=f"Epoch [{epoch+1}/{num_epochs_fc}]")
38     for images, labels in loop:
39         images, labels = images.to(device), labels.to(device)
40         fc_optimizer.zero_grad() # Zero the gradients
41         outputs = model_genre(images) # Forward pass
42         loss = criterion(outputs, labels) # Compute loss
43         loss.backward() # Backward pass to compute gradients
44         fc_optimizer.step() # Update FC layer weights
45
46         running_loss += loss.item()
47         _, predicted = torch.max(outputs, 1)
48         correct += (predicted == labels).sum().item()
49         total += labels.size(0)
50     loop.set_postfix(loss=running_loss / len(genre_train_loader), acc=100 * correct / total)
51
52     print(f"Epoch [{epoch+1}/{num_epochs_fc}], Loss: {running_loss / len(genre_train_loader):.4f}, "
53           f"Accuracy: {100 * correct / total:.2f}%")
54
55     print("🔴 Step 1 Complete: FC Layer Training Finished!")
56
57 # Step 2: Fine-Tuning Entire Network
58 def fine_tune():
59     print("Starting Step 2: Fine-Tuning the Entire Network...")
60     for param in model_genre.parameters():
61         param.requires_grad = True # Unfreeze all layers
62
63     # Define optimizer and scheduler for fine-tuning
64     finetune_optimizer = optim.Adam(model_genre.parameters(), lr=finetune_lr, weight_decay=weight_decay)
65     scheduler = ReduceLROnPlateau(finetune_optimizer, mode='min', factor=0.1, patience=3, verbose=True)
66
67     best_val_loss = float('inf')
68
69     for epoch in range(num_epochs_finetune):
70         model_genre.train()
71         running_loss, correct, total = 0.0, 0, 0
72         loop = tqdm(genre_train_loader, leave=True, desc=f"Epoch [{epoch+1}/{num_epochs_finetune}]")
73
74         for images, labels in loop:
75             images, labels = images.to(device), labels.to(device)
76             finetune_optimizer.zero_grad()
77             outputs = model_genre(images)
78             loss = criterion(outputs, labels)
79             loss.backward()
80             finetune_optimizer.step()
81
82             running_loss += loss.item()
83             _, predicted = torch.max(outputs, 1)
84             correct += (predicted == labels).sum().item()
85             total += labels.size(0)
86             loop.set_postfix(loss=running_loss / len(genre_train_loader), acc=100 * correct / total)
87
88     # Validation phase
89     val_loss, val_acc = evaluate(gnere_val_loader)

```

```

90     scheduler.step(val_loss)
91
92     print(f"Epoch [{epoch+1}/{num_epochs_finetune}], "
93           f"Train Loss: {running_loss / len(genre_train_loader):.4f}, Train Acc: {100 * correct / total:.2f}%, "
94           f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")
95
96     # Save the best model based on validation loss
97     if val_loss < best_val_loss:
98         best_val_loss = val_loss
99         torch.save(model_genre.state_dict(), "/content/drive/MyDrive/art_model/best_resnet18_genre_model_450.pth")
100        print("✅ Best model saved!")
101
102    print("⌚ Step 2 Complete: Fine-Tuning Finished!")
103
104 # Function to evaluate the model on validation data
105 def evaluate(loader):
106     model_genre.eval()
107     val_loss, correct, total = 0.0, 0, 0
108
109     with torch.no_grad():
110         for images, labels in loader:
111             images, labels, _ = images.to(device), labels.to(device)
112             outputs = model_genre(images)
113             loss = criterion(outputs, labels)
114             val_loss += loss.item()
115             _, predicted = torch.max(outputs, 1)
116             correct += (predicted == labels).sum().item()
117             total += labels.size(0)
118
119     avg_loss = val_loss / len(loader)
120     accuracy = 100 * correct / total
121     return avg_loss, accuracy
122
123
124 # Start Training: Step 1 (Train FC Layer) and Step 2 (Fine-Tuning)
125 train_fc_layer() # Train the FC layer only
126 fine_tune() # Fine-tune the entire ResNet model
127
128 print("🎉 Training complete! The best model is saved as 'best_resnet18_model.pth'.")
```

⌚ Starting Step 1: Training FC Layer Only...

Epoch [1/5]: 100%|██████████| 141/141 [10:06<00:00, 4.30s/it, acc=46.3, loss=1.63]
 Epoch [1/5], Loss: 1.6341, Accuracy: 46.31%
 Epoch [2/5]: 100%|██████████| 141/141 [01:26<00:00, 1.63it/s, acc=62.6, loss=1.18]
 Epoch [2/5], Loss: 1.1809, Accuracy: 62.56%
 Epoch [3/5]: 100%|██████████| 141/141 [01:28<00:00, 1.59it/s, acc=65.3, loss=1.08]
 Epoch [3/5], Loss: 1.0802, Accuracy: 65.27%
 Epoch [4/5]: 100%|██████████| 141/141 [01:29<00:00, 1.57it/s, acc=65.8, loss=1.03]
 Epoch [4/5], Loss: 1.0342, Accuracy: 65.80%
 Epoch [5/5]: 100%|██████████| 141/141 [01:24<00:00, 1.66it/s, acc=66.5, loss=0.993]
 /usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_warnings.warn()
 Epoch [5/5], Loss: 0.9928, Accuracy: 66.53%

⌚ Step 1 Complete: FC Layer Training Finished!

Starting Step 2: Fine-Tuning the Entire Network...

Epoch [1/5]: 100%|██████████| 141/141 [01:28<00:00, 1.59it/s, acc=68.4, loss=0.96]
 Epoch [1/5], Train Loss: 0.9598, Train Acc: 68.44%, Val Loss: 0.9048, Val Acc: 71.50%
 ✅ Best model saved!
 Epoch [2/5]: 100%|██████████| 141/141 [01:27<00:00, 1.61it/s, acc=92.8, loss=0.277]
 Epoch [2/5], Train Loss: 0.2768, Train Acc: 92.78%, Val Loss: 0.8774, Val Acc: 71.50%
 ✅ Best model saved!
 Epoch [3/5]: 100%|██████████| 141/141 [01:33<00:00, 1.51it/s, acc=98.6, loss=0.0765]
 Epoch [3/5], Train Loss: 0.0765, Train Acc: 98.64%, Val Loss: 0.9237, Val Acc: 72.80%
 Epoch [4/5]: 100%|██████████| 141/141 [01:28<00:00, 1.60it/s, acc=99.9, loss=0.0291]
 Epoch [4/5], Train Loss: 0.0291, Train Acc: 99.89%, Val Loss: 0.9261, Val Acc: 73.40%
 Epoch [5/5]: 100%|██████████| 141/141 [01:28<00:00, 1.59it/s, acc=99.9, loss=0.0156]
 Epoch [5/5], Train Loss: 0.0156, Train Acc: 99.91%, Val Loss: 0.9837, Val Acc: 73.10%

⌚ Step 2 Complete: Fine-Tuning Finished!

🎉 Training complete! The best model is saved as 'best_resnet18_model.pth'.

1 len(genre_class)

⌚ 10

1 len(artist_balanced_dataset), len(style_balanced_dataset), len(genre_balanced_dataset)

```
Traceback (most recent call last)
<ipython-input-19-b78109855c23> in <cell line: 0>()
----> 1 len(artist_balanced_dataset), len(style_balanced_dataset), len(genre_balanced_dataset)

NameError: name 'style_balanced_dataset' is not defined
```

1