

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1
2 import pandas as pd
3 csv_path = './content/MyDrive/drive/data/csv'
4 artist_train = pd.read_csv('/content/drive/MyDrive/data/csv/Artist/artist_train')
5 # lets visualize one imag
6 base_url = '/content/drive/MyDrive/data/images'
7 # lets start creating data
8 artist = '/content/drive/MyDrive/data/csv/Artist'
9 genre = '/content/drive/MyDrive/data/csv/Genre'
10 style = '/content/drive/MyDrive/data/csv/Style'
11 data_dir = '/content/drive/MyDrive/data/csv'
12
13 artist_train_path = data_dir + '/artist_train.csv'
14 artist_val_path = data_dir + '/artist_val.csv'
15 artist_class_path = data_dir + '/artist_class.txt'
16
17 genre_train_path = data_dir + '/genre_train.csv'
18 genre_val_path = data_dir + '/genre_val.csv'
19 genre_class_path = data_dir + '/genre_class.txt'
20
21 style_train_path = data_dir + '/style_train.csv'
22 style_val_path = data_dir + '/style_val.csv'
23 style_class_path = data_dir + '/style_class.txt'
24
25 artist_train = pd.read_csv(data_dir + '/artist_train.csv')
26 artist_val = pd.read_csv(data_dir + '/artist_val.csv')
27 artist_class = pd.read_csv(artist_class_path, header=None, names=["artist_name"])
28
29 genre_train = pd.read_csv(data_dir + '/genre_train.csv')
30 genre_val = pd.read_csv(data_dir + '/genre_val.csv')
31 genre_class = pd.read_csv(genre_class_path, header=None, names=["genre_name"])
32
33
34 style_train = pd.read_csv(data_dir + '/style_train.csv')
35 style_val = pd.read_csv(data_dir + '/style_val.csv')
36 style_class = pd.read_csv(style_class_path, header=None, names=["style_name"])
37
38 # genre_class['genre_name'][1]
39 len(style_class)
```

27

```
1 import os
2 import pandas as pd
3 import torch
4 from torch.utils.data import Dataset
5 from torchvision import transforms
6 from PIL import Image
7 from collections import defaultdict
8 import random
9 from tqdm import tqdm
10 import matplotlib.pyplot as plt
11
12 # Define dataset class
13 class BalancedArtDataset(Dataset):
14     def __init__(self, csv_file, img_dir, class_mapping, transform=None, images_per_class=32):
15         self.data = pd.read_csv(csv_file)
16         self.img_dir = img_dir
17         self.class_mapping = class_mapping
18         self.transform = transform
19         self.images_per_class = images_per_class
20
21         # Filter out missing images
22         print("Filtering missing images...")
23         self.data = self.data[self.data.iloc[:, 0].apply(lambda x: os.path.exists(os.path.join(img_dir, str(x))))]
24
25         # Group images by class
26         print("Grouping images by class...")
27         self.class_images = defaultdict(list)
```

```

28     for _, row in tqdm(self.data.iterrows(), total=len(self.data), desc="Processing rows"):
29         self.class_images[row.iloc[1]].append(row)
30
31     # Balance dataset with 32 images per class
32     print("Balancing dataset...")
33     self.final_data = []
34     all_images = []
35     for cls, images in tqdm(self.class_images.items(), total=len(self.class_images), desc="Processing classes"):
36         if len(images) >= images_per_class:
37             selected_images = random.sample(images, images_per_class)
38         else:
39             selected_images = images[:]
40             all_images.extend(images) # Store extra images for filling
41             self.final_data.extend(selected_images)
42
43     # Fill missing slots with extra images
44     print("Filling missing slots...")
45     needed_images = images_per_class * len(self.class_images) - len(self.final_data)
46     if needed_images > 0:
47         self.final_data.extend(random.sample(all_images, min(needed_images, len(all_images))))
48
49     # Shuffle dataset
50     print("Shuffling dataset...")
51     random.shuffle(self.final_data)
52
53     # Count images per class
54     self.class_counts = defaultdict(int)
55     for row in self.final_data:
56         self.class_counts[row.iloc[1]] += 1
57
58     def __len__(self):
59         return len(self.final_data)
60
61     def __getitem__(self, idx):
62         row = self.final_data[idx]
63         img_path = os.path.join(self.img_dir, str(row.iloc[0]))
64         label = row.iloc[1]
65         image = Image.open(img_path).convert("RGB")
66
67         if self.transform:
68             image = self.transform(image)
69
70         return image, label
71
72
73     def visualize_class_distribution(self):
74         plt.figure(figsize=(12, 6))
75         plt.bar(self.class_counts.keys(), self.class_counts.values(), color='skyblue')
76         plt.xlabel("Class")
77         plt.ylabel("Number of Images")
78         plt.title("Class Distribution in Balanced Dataset")
79         plt.xticks(rotation=45)
80         plt.show()
81
82     def visualize_samples(self, num_samples=10):
83         fig, axes = plt.subplots(1, num_samples, figsize=(40, 20))
84         for i in range(num_samples):
85             image, label = self.__getitem__(random.randint(0, len(self) - 1))
86             image = image.permute(1, 2, 0).numpy() # Convert to (H, W, C)
87             image = (image * 0.5) + 0.5 # Unnormalize
88             axes[i].imshow(image)
89             axes[i].set_title(f"Class: {label} {genre_class['genre_name'][label]}")
90             axes[i].axis("off")
91         plt.show()
92
93     # Function to compare artist and genre relationships
94
95
96
97 # Define transformations
98 transform = transforms.Compose([
99     transforms.Resize((224, 224)),
100    transforms.ToTensor(),
101    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
102 ])
103
104 # Create balanced artist dataset

```

```


105 print("Creating balanced artist dataset...")
106 genre_balanced_dataset = BalancedArtDataset(genre_train_path, "/content/drive/MyDrive/data/images", genre_class_path, transform=transfc
107 genre_test_balanced_dataset = BalancedArtDataset(genre_val_path, "/content/drive/MyDrive/data/images", genre_class_path, transform=tr
108
109 # Check dataset length
110 print("Artist balanced dataset size:", len(genre_balanced_dataset), len(genre_test_balanced_dataset))
111
112 # Visualize class distribution
113 genre_balanced_dataset.visualize_class_distribution()
114
115 # Visualize sample images
116 genre_balanced_dataset.visualize_samples(num_samples=10)
117

```

```

↗ Creating balanced artist dataset...
Filtering missing images...
Grouping images by class...
Processing rows: 100%|██████████| 45501/45501 [00:03<00:00, 13441.81it/s]
Balancing dataset...
Processing classes: 100%|██████████| 10/10 [00:00<00:00, 2948.75it/s]
Filling missing slots...
Shuffling dataset...
Filtering missing images...
Grouping images by class...
Processing rows: 100%|██████████| 19492/19492 [00:01<00:00, 19431.19it/s]
Balancing dataset...
Processing classes: 100%|██████████| 10/10 [00:00<00:00, 5752.71it/s]
Filling missing slots...
Shuffling dataset...
Artist balanced dataset size: 4500 1500

```



```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision.models as models
5 from torch.utils.data import DataLoader
6 from tqdm import tqdm
7
8 # Check for GPU
9 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10 print(f"Using device: {device}")
11
12 # Define DataLoaders
13 batch_size = 32 # Adjust based on GPU memory
14
15 train_loader = DataLoader(genre_balanced_dataset, batch_size=batch_size, shuffle=True, num_workers=4, pin_memory=True)
16 val_loader = DataLoader(genre_test_balanced_dataset, batch_size=batch_size, shuffle=False, num_workers=4, pin_memory=True)
17
18 print("DataLoaders created successfully!")
19

```

```

↗ Using device: cuda
DataLoaders created successfully!
/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes
warnings.warn(

```

```

1 # Import necessary libraries
2 import torch
3 import torch.nn as nn
4 import torch.optim as optim
5 from torch.optim.lr_scheduler import ReduceLROnPlateau
6 from torchvision import datasets, transforms, models
7 from torch.utils.data import DataLoader
8 from tqdm import tqdm
9 import time
10
11 # Hyperparameters
12 num_epochs_fc = 5          # Train FC layer only for 5 epochs
13 num_epochs_finetune = 6   # Fine-tune entire model for 6 more epochs
14 initial_lr = 0.001        # Learning rate for training FC layer
15 finetune_lr = 1e-4        # Lower learning rate for fine-tuning
16 batch_size = 32
17 weight_decay = 1e-4       # L2 regularization
18
19 # Define the number of output classes
20 num_classes = len(genre_class) # Replace with your actual class count
21
22 # Load the ResNet-50 model pre-trained on ImageNet
23 model = models.resnet50(weights=models.ResNet50_Weights.DEFAULT)
24 model.fc = nn.Linear(model.fc.in_features, num_classes).to(device) # Replace
    the FC layer to match your task
25
26 model = model.to(device) # Move model to the appropriate device (CPU/GPU)
27
28 # Define loss function and optimizer for training the FC layer
29 criterion = nn.CrossEntropyLoss().to(device)
30 fc_optimizer = optim.Adam(model.fc.parameters(), lr=initial_lr,
    weight_decay=weight_decay)
31
32 # Function to train the fully connected (FC) layer only
33 def train_fc_layer():
34     print("Starting Step 1: Training FC Layer Only...")
35
36     # Freeze all layers except the FC layer
37     for param in model.parameters():
38         param.requires_grad = False
39
40     # Ensure the new FC layer is trainable
41     for param in model.fc.parameters():
42         param.requires_grad = True
43
44     model.train() # Set the model to training mode
45
46     for epoch in range(num_epochs_fc):
47         running_loss, correct, total = 0.0, 0, 0
48         loop = tqdm(train_loader, leave=True, desc=f"Epoch [{epoch+1}]/
            {num_epochs_fc}")
49
50         for images, labels in loop:
51             images, labels = images.to(device), labels.to(device)
52
53             # Zero the gradients
54             fc_optimizer.zero_grad()
55
56             # Forward pass
57             outputs = model(images)
58
59             # Compute loss
60             loss = criterion(outputs, labels)
61             loss.backward() # Backpropagation
62             fc_optimizer.step() # Update FC layer weights
63
64             running_loss += loss.item()
65             _, predicted = torch.max(outputs, 1)
66             correct += (predicted == labels).sum().item()
67             total += labels.size(0)
68
69         # Update the progress bar
70         loop.set_postfix(loss=running_loss / len(train_loader), acc=100 *
            correct / total)
71
72     print(f"Epoch [{epoch+1}]/{num_epochs_fc}], Loss: {running_loss / len
    (train_loader): %.4f" %

```

```

    (train_loader).4T},
73         f"Accuracy: {100 * correct / total:.2f}%")
74
75     print("🔴 Step 1 Complete: FC Layer Training Finished!")
76
77
78 # Function to fine-tune the entire ResNet-50 network
79 def fine_tune():
80     print("Starting Step 2: Fine-Tuning the Entire Network...")
81
82     # Unfreeze all layers
83     for param in model.parameters():
84         param.requires_grad = True
85
86     # Define optimizer and learning rate scheduler for fine-tuning
87     finetune_optimizer = optim.Adam(model.parameters(), lr=finetune_lr,
88     weight_decay=weight_decay)
89     scheduler = ReduceLROnPlateau(finetune_optimizer, mode='min', factor=0.1,
90     patience=3, verbose=True)
91
92     best_val_loss = float('inf')
93
94     for epoch in range(num_epochs_finetune):
95         model.train()
96         running_loss, correct, total = 0.0, 0, 0
97         loop = tqdm(train_loader, leave=True, desc=f"Epoch [{epoch+1}/
98         {num_epochs_finetune}]")
99
100        for images, labels in loop:
101            images, labels = images.to(device), labels.to(device)
102
103            finetune_optimizer.zero_grad()
104
105            # Forward pass
106            outputs = model(images)
107
108            # Compute loss
109            loss = criterion(outputs, labels)
110            loss.backward()
111            finetune_optimizer.step()
112
113            running_loss += loss.item()
114            _, predicted = torch.max(outputs, 1)
115            correct += (predicted == labels).sum().item()
116            total += labels.size(0)
117
118            # Update the progress bar
119            loop.set_postfix(loss=running_loss / len(train_loader), acc=100 *
120            correct / total)
121
122        # Validation phase
123        val_loss, val_acc = evaluate(val_loader)
124        scheduler.step(val_loss)
125
126        print(f"Epoch [{epoch+1}/{num_epochs_finetune}], "
127        f"Train Loss: {running_loss / len(train_loader):.4f}, Train Acc:
128        {100 * correct / total:.2f}%, "
129        f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%")
130
131        # Save the best model based on validation loss
132        if val_loss < best_val_loss:
133            best_val_loss = val_loss
134            torch.save(model.state_dict(), "/content/drive/MyDrive/art_model/
135            best_genre_resnet50_model.pth")
136            print("✅ Best model saved!")
137
138    print("🔴 Step 2 Complete: Fine-Tuning Finished!")
139
140
141 # Function to evaluate the model on validation data
142 def evaluate(loader):
143     model.eval()
144     val_loss, correct, total = 0.0, 0, 0
145
146     with torch.no_grad():
147         for images, labels in loader:
148             images, labels = images.to(device), labels.to(device)
149
150

```

```

144         outputs = model(images)
145         loss = criterion(outputs, labels)
146         val_loss += loss.item()
147         _, predicted = torch.max(outputs, 1)
148         correct += (predicted == labels).sum().item()
149         total += labels.size(0)
150
151     avg_loss = val_loss / len(loader)
152     accuracy = 100 * correct / total
153     return avg_loss, accuracy
154
155
156 # Start training
157 train_fc_layer() # Step 1: Train FC Layer only
158 fine_tune()      # Step 2: Fine-tune the entire ResNet-50 model
159
160 print("🎉 Training complete! The best model is saved as 'best_resnet50_model.
    pth'.")
161

```

Starting Step 1: Training FC Layer Only...

```

Epoch [1/5]: 100%|██████████| 141/141 [07:49<00:00, 3.33s/it, acc=55.4, loss=1.45]
Epoch [1/5], Loss: 1.4543, Accuracy: 55.36%
Epoch [2/5]: 100%|██████████| 141/141 [01:16<00:00, 1.83it/s, acc=68.2, loss=1.01]
Epoch [2/5], Loss: 1.0138, Accuracy: 68.24%
Epoch [3/5]: 100%|██████████| 141/141 [01:18<00:00, 1.81it/s, acc=71.8, loss=0.905]
Epoch [3/5], Loss: 0.9046, Accuracy: 71.82%
Epoch [4/5]: 100%|██████████| 141/141 [01:17<00:00, 1.83it/s, acc=74, loss=0.819]
Epoch [4/5], Loss: 0.8189, Accuracy: 74.02%
Epoch [5/5]: 100%|██████████| 141/141 [01:19<00:00, 1.78it/s, acc=75.6, loss=0.775]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get
warnings.warn(
Epoch [5/5], Loss: 0.7755, Accuracy: 75.58%
🎉 Step 1 Complete: FC Layer Training Finished!
Starting Step 2: Fine-Tuning the Entire Network...
Epoch [1/6]: 100%|██████████| 141/141 [01:27<00:00, 1.61it/s, acc=76.9, loss=0.678]
Epoch [1/6], Train Loss: 0.6778, Train Acc: 76.89%, Val Loss: 0.8499, Val Acc: 72.67%
✅ Best model saved!
Epoch [2/6]: 100%|██████████| 141/141 [01:28<00:00, 1.58it/s, acc=95.8, loss=0.181]
Epoch [2/6], Train Loss: 0.1805, Train Acc: 95.78%, Val Loss: 0.9826, Val Acc: 70.47%
Epoch [3/6]: 100%|██████████| 141/141 [01:27<00:00, 1.62it/s, acc=99.4, loss=0.0472]
Epoch [3/6], Train Loss: 0.0472, Train Acc: 99.44%, Val Loss: 1.0296, Val Acc: 71.73%
Epoch [4/6]: 100%|██████████| 141/141 [01:25<00:00, 1.65it/s, acc=99.6, loss=0.0241]
Epoch [4/6], Train Loss: 0.0241, Train Acc: 99.64%, Val Loss: 1.1336, Val Acc: 72.13%
Epoch [5/6]: 100%|██████████| 141/141 [01:27<00:00, 1.61it/s, acc=100, loss=0.0134]
Epoch [5/6], Train Loss: 0.0134, Train Acc: 99.96%, Val Loss: 1.1553, Val Acc: 72.67%
Epoch [6/6]: 100%|██████████| 141/141 [01:27<00:00, 1.62it/s, acc=100, loss=0.00832]
Epoch [6/6], Train Loss: 0.0083, Train Acc: 99.96%, Val Loss: 1.1345, Val Acc: 72.67%
🎉 Step 2 Complete: Fine-Tuning Finished!
🎉 Training complete! The best model is saved as 'best_resnet50_model.pth'.

```